

# Ensemble Learning and Random Forests

## Ensembles

- If we aggregate the predictions of a group of predictors, we will often get better predictions than with the best individual predictor.
- A group of predictors is called an ensemble, and aggregating individual predictors is called ensemble learning. An ensemble learning algorithm is called an ensemble method.
- For example, we can train a group of decision tree classifiers, each on a different subset of the training set, get their predictions, and get the class with the most votes. This ensemble is called a random forest.

## Voting Classifiers

- We can aggregate the predictions of each classifier and get the majority vote. This classifier is a hard voting classifier.
- Even if each classifier is a weak learner, the ensemble can still be a strong learner.
  - There has to be a sufficient number of weak learners in the ensemble, and they must be sufficiently diverse to create a strong learner.
  - One way to get diverse learners is to train them using different algorithms.
- We can use Scikit-Learn's **VotingClassifier()** class, which takes in a list of name/predictor tuples.
  - Note that moving forward, every class or function mentioned originates from the Scikit-Learn API unless specified otherwise.
- The "estimators" attribute fits a clone of every predictor that was provided to it.
- Then, we can simply fit() and predict().
- Voting classifiers can also perform soft voting, where the algorithm predicts the highest class probability, averaged over all individual classifiers.
  - Soft voting often performs better than hard voting because it gives more weight to highly confident votes.
  - We simply set the voting parameter to "soft."

## Bagging and Pasting

- Another way of getting diverse classifiers is to use the same training algorithm for every predictor but train them on different random subsets of the training set.

- Bagging is sampling performed with replacement, also known as bootstrap aggregating.
  - Pasting is sampling performed without replacement.
- The aggregation function is typically the statistical mode for classification (i.e., the most frequent prediction, just like with a hard voting classifier) or the average for regression.
- We can use the **BaggingClassifier()** class (or **BaggingRegressor()** class for regression) with a **DecisionTreeClassifier()** as an input.
  - By default, this class performs bootstrapping. To perform pasting, we set the “bootstrap” parameter to false.
  - The bagging classifier also performs soft voting by default.
- An ensemble’s predictions likely generalize much better than a single decision tree.
  - An ensemble has a similar amount of bias but a smaller variance.
  - Bagging introduces more bias than pasting, but the overall ensemble’s variance is smaller.

### Out-of-Bag Evaluation

- Not all training instances are sampled for any given predictor. The instances that are not sampled are called out-of-bag (OOB) instances.
- A bagging ensemble can be evaluated using OOB instances without a separate validation set.
- We can simply set the bagging classifier’s “oob\_score” attribute to true and check its evaluation score with the **oob\_score\_** attribute.
- The classifier also has an **oob\_decision\_function\_** attribute to get the class probabilities for each instance.

### Random Patches and Random Subspaces

- The bagging classifier can also sample features.
- We can use the “bootstrap\_features” hyperparameter to train each predictor on a random subset of input features.
- Sampling both instances and features is called the random patches method.
- Keeping all training instances except the sampling features is called the random subspaces method.

## Random Forests

- We can use **RandomForestClassifier()** to use an ensemble of decision trees.
  - There’s also the **RandomForestRegressor()** class for regression tasks.
- The random forest algorithm searches for the best feature among a random subset of features.
  - It samples  $\sqrt{n}$  features, where  $n$  is the total number of features.
  - This results in a higher bias but lower variance.
- The **RandomForestClassifier()** class has the same hyperparameters as a **DecisionTreeClassifier()** and as the **BaggingClassifier()** class.

- We can find the most important features using the RandomForestClassifier's **feature\_importances\_** attribute.

### Extra-Trees

- We can make trees even more random by using random thresholds for each feature.
  - We simply set the “splitter” attribute to “random” when creating a decision tree classifier.
  - These forests are called extremely randomized trees or extra-trees.
- Extra-Trees have a higher bias but lower variance than random forests.
- Extra-Tree classifiers are much faster than random forests because it simply uses a random threshold at each feature rather than finding the best possible threshold at each feature.
- We can use the **ExtraTreesClassifier()** class whose API is identical to the RandomForestClassifier() class.

### Boosting

- Boosting refers to any ensemble method that can combine weak learners into a strong learner.
- Boosting methods train predictors sequentially, each trying to correct its predecessor.

### AdaBoost

- An AdaBoost classifier pays more attention to the training instances that the predecessor underfit.
  - The algorithm first trains a base classifier, and then, it increases the relative weight of misclassified training instances.
  - The algorithm repeats this process a set number of times.
- We can use the **AdaBoostClassifier()** class using a base classifier, like a decision tree, as an input.
- An important drawback of AdaBoost is that it does not scale as well as bagging or pasting.
- If the AdaBoost ensemble overfits, we can try to reduce the number of estimators or regularizing the base estimator more.

### Gradient Boosting

- This method tries to fit the new predictor to the residual errors of the predecessor rather than the relative weight of misclassified training instances.
- We can use the **GradientBoostingRegressor()** class or the **GradientBoostingClassifier()** class depending on the task.
  - These classes have a “learning\_rate” hyperparameter that controls the contribution of each decision tree.
  - The the learning rate value is lower, then the algorithm needs more trees in the ensemble, but it will likely generalize better. This technique is called shrinkage.

- If the learning rate value is too low, the model will likely underfit, and if it is set too high, it will likely overfit.
- To find the optimal number of trees, we can perform cross-validation using `GridSearchCV()` or `RandomizedSearchCV()`.
  - We can also set the gradient booster class's "n\_iter\_no\_change" attribute to an integer value, which will stop adding more trees during training if the last few trees did not help.
  - Then, we can simply access the optimal number of trees using the class's **n\_estimators\_** attribute.

### Histogram-Based Gradient Boosting

- HGB bins the input features, replacing them with integers.
  - The number of bins is controlled by the "max\_bins" hyperparameter, which defaults to 255 and cannot go any higher.
  - Working with integers allows the algorithm to use faster and more memory-efficient data structures.
  - This implementation has a complexity of  $O(b \times m)$  where  $b$  is the number of bins and  $m$  is the number of training instances.
- We can use the **HistGradientBoostingRegressor()** class or the **HistGradientBoostingClassifier()** class.
  - Early stopping is automatically activated if the number of instances is greater than 10,000. This is controlled by setting the "early\_stopping" hyperparameter to true or false.
- HGB classes support categorical features and missing values, which means that we can use an encoder to encode text as integers.

### Stacking

- Stacked generalization trains a model to perform aggregation.
  - Each predictor predicts a different value, and the final predictor (the blender or meta learner) takes these predictions as inputs and makes the final prediction.
- The blending training set will contain one input feature per predictor.
- Once the algorithm trains the blender, it trains the base predictors one last time on the full original training set.
- We can have several layers, or stacks, of blenders, but this increases both training time and system complexity.
- We can use the **StackingClassifier()** class, which has a similar API as the `VotingClassifier()` class.
  - The `StackingClassifier()` class has a "final\_estimator" attribute, which controls which estimator, such as a random forest, the algorithm uses as the blender.